

Assessing Practitioner Beliefs

N.C. Shrikanth, Tim Menzies
snaraya7@ncsu.edu, timm@ieee.org
North Carolina State University, USA

ABSTRACT

Just because software developers say they believe in “X” that does not necessarily mean that “X” is true. As shown here, there exist numerous beliefs listed in the recent Software Engineering literature which are only supported by small portions of the available data. Hence we ask what is the source of this disconnect between beliefs and evidence?.

To answer this question we look for evidence for ten beliefs within 300,000+ changes seen in dozens of open-source projects. Some of those beliefs had strong support across all the projects; specifically, “A commit that involves more added and removed lines is more bug-prone” and “Files with fewer lines contributed by their owners (who contribute most changes) are bug-prone”.

Most of the widely-held beliefs studied are only sporadically supported in the data; i.e. large effects can appear in project data and then disappear in subsequent releases. Such sporadic support explains why developers believe things that were relevant to their prior work, but not necessarily their current work.

Our conclusion will be that we need to change the nature of the debate with SE. Specifically, while it is important to report the effects that hold right now, it is also important to report on what effects change over time.

KEYWORDS

defects, beliefs, practitioner, empirical software engineering

1 INTRODUCTION

Just because software developers say they believe in “X”, that does not necessarily mean that “X” is true. Jørgensen & Gruschke [12] note that in software engineering domain, seldom lessons from past projects are used to improve future reasoning (to the detriment of new projects). Passos et al. [26] note that developers often assume that lessons learned from a few past projects are general to all future projects [26]. Devanbu et al. record opinions about software development from 564 Microsoft software developers from around the world [4]. They comment that programmer beliefs can (a) vary with each project and (b) may not necessarily correspond with actual evidence in their current projects. Menzies and Nagappan et al. offer specific examples of this effect. Nagappan et al. [24] have shown that the much-feared *goto* statement is usually benign, and sometimes even useful. Menzies et al. recently found no evidence for the *delayed issue effect* [21] (“*The longer a bug remains in the system, the exponentially more costly it becomes to fix*”) in 171 industry projects. That study shows that a widely held belief by both practitioners and academics cannot be assumed to always hold.

Accordingly, we think it is important to always carefully evaluate developer beliefs. For example, Table 1 lists ten beliefs and software defects listed in a recent TSE’18 paper by Wan et al. [32]. That study

collected 395 responses from practitioners to document developer beliefs about willingness to adopt technologies, challenges, defect prediction metrics, etc. This paper looks for evidence for the Table 1 beliefs in 300,000+ changes seen in dozens of open-source projects. What we find is:

- Two of these beliefs are supported in the data; specifically, “A commit that involves more added and removed lines is more bug-prone” and “Files with fewer lines contributed by their owners (who contribute most changes) are bug-prone”.
- The beliefs supported by the data are not usually endorsed by developers. The last column of Table 1 shows the percent of the 395 developers who endorsed each belief. Note that the last belief (B10) is only endorsed by 30% of the surveyed developers. Yet in our data, it is one of the strongest effects.
- As to the other eight items listed in Table 1, while these have sporadic support in different releases of the same project, there is little overall evidence for these beliefs.

Our conclusion will be that we need to change the nature of the debate with SE. While it is important to report the effects that hold right now, it is also important to report on what effects change over time.

The rest of this paper is structured as follows. §2 relates the current work to the prior literature on metrics and instability. In §3, we discuss the choice of our datasets, its distribution and brief about our statistical methods we use. §4 details the construction of experiments mapping to literature. Results in §5 and directions to future work in §6. §7 discusses the reliability of our findings. Finally, the conclusions are presented in §8.

This paper reports strongly supported beliefs and discusses prevalence of beliefs. To make that argument we will use the list of developer beliefs about defect prediction in Table 1. Using that list, we will ask:

Table 1: Practitioners’ agreement % on defect prediction metrics (beliefs) reported in a recent IEEE Transactions on Software Engineering paper by Wan et al. [32].

#	Belief	%
B1	A file with a complex code change process tends to be buggy.	76
B2	A file that is changed by more developers is more bug-prone.	64
B3	A file with more added lines is more bug-prone.	61
B4	Recently changed files tend to be buggy.	58
B5	A commit that involves more added and removed lines is more bug-prone.	57
B6	Recently bug-fixed files tend to be buggy	49
B7	A file with more fixed bugs tends to be more bugprone.	48
B8	A file with more commits is more bug-prone.	46
B9	A file with more removed lines is more bug-prone.	35
B10	Files with fewer lines contributed by their owners (who contribute most changes) are bug-prone.	30

RQ1 : What beliefs are strongly supported in the data?**Result:**

We found that beliefs labeled B10 & B5 in Table 1, which is believed by only (30% & 57%) of practitioners has strong support (i.e., large correlations), whereas the more popular belief B1, which is believed by 76% of practitioners, has relatively weaker support.

And to identify the source of disconnect between beliefs and evidence, next we ask:

RQ2 : Does the same evidence appear everywhere?

- Do projects show evidence for all the beliefs?
- Does the size of a release affect belief support?
- Do beliefs evolve as a project matures (more releases)?

**Result:**

Only 24% of the projects show support for all the 10 beliefs. And those projects showed support for beliefs among 10% - 36% of their releases. Beliefs appear stronger in smaller than larger releases, with fluctuations like B4 (weaker in *large* releases but stronger in *small* releases). Support for beliefs tended to decay than strengthen as the project matured.

2 BACKGROUND

2.1 Why Study Practitioner Beliefs?

If we do not understand what factors lead to software defects, then that has detrimental effects for quality assurance, trust, insight, training, and tool development.

- *Quality assurance*: If we do not know what causes defects in a project, we cannot prevent those problems.
- *Training*: Another concern is what do we train novice software engineers or newcomers to a project? If our models are not stable, then it is hard to teach what factors most influence software quality.
- *Tool development*: Further to the last point— if we are unsure what factors most influence quality, it is difficult to design, implement and deploy tools that can successfully improve that quality (e.g. Static analysis tools).

A premise of much research [16, 32–34, 37] is that knowledge of what factors influence software quality can be discovered by asking software practitioners. But based on the results of this paper, we would say that:

- While such surveys produce a long list of possible effects,
- Only some of which might be relevant to a particular project at a particular time.

Hence, we say:

- Yes, we should study working software engineers to learn a list of effects that might damage software quality;
- But no, do not assume that all those effects hold at all times over the current project.

More specifically, when setting policies for software projects, managers and project leads should not “bet” on a small number of effects. Rather, they should:

- Perpetually monitor for the presence, or absence of a range of effects (such as those listed in Table 1);
- Perpetually adjust their code review and code refactoring processes such that they learn (a) not to trigger on old effects that now no longer hold or (b) trigger on new effects that have just appeared.

2.2 Related Work

Perhaps the closest work to our work is the work by McIntosh & Kamei who, in 2017, reported that the effect on defects of code change properties based on size, entropy, history, and code ownership changes as software systems evolve [19]. Further, they report that although size-related metrics (B3:*Added lines* & B9:*Removed lines*) were their top contributors over other metrics, that effect tended to fluctuate across time (and the strength of that effect was project-specific).

To start the research of this paper, we found five papers [16, 32–34, 37] which list practitioner beliefs, but do not test those beliefs with respect to empirical evidence. From this, we chose a recent qualitative study for IEEE Transactions of Software (see Table 1). This paper performed a comprehensive study by gathering practitioners’ opinions (395 responses) on various defect prediction research hypotheses discussed in the literature between 2012-2017. More importantly, they highlight practitioners’ agreement % on defect prediction metrics, prioritization strategy, etc. In support of that work, we work that many of the beliefs reported in [32] by Wan et al. were also reported previously in [3, 13]. In this paper, we revisit these beliefs independently to look for current evidence and reasoning disconnect between beliefs and evidence using the prevalence of their support.

3 METHODOLOGY OVERVIEW

In this section, we elucidate source, distribution and various attributes we collect from our sample projects. Then we detail the underlying statistical techniques we use to assess our beliefs.

3.1 Data Source & Cleaning

For this work, we wanted to study conclusion instability using a much more detailed approach than the *Related Work* mentioned above. We analyze 3 times more changes (commits) than recent defect prediction work [11, 19] and the volume of our dataset is 8 times larger as we expand those changes (commits) that results in 301,627 source code file entries filtered from 524,851 in total.

Turning then to Github¹, we chose 50 GitHub repository links from a recent defect prediction paper [31] that uses Munaiah et al.’s GitHub project database [23]. Munaiah et al.’s work simplified choosing substantive samples from GitHub attributing to several metrics such as maturity, active developers, license, etc., We mined file-level commit histories and release information using Github API. To minimize bias in our modelling we applied the following sanity checks and discarded projects that have less than,

¹<https://github.com>

Table 2: 37 OS projects developed in popular programming languages. Some projects are developed using multiple programming languages. (Rb - Ruby, Py - Python, JS - JavaScript and SCSS - Sassy-CSS)

URL (github.com/)	Language	URL (github.com/)	Language
activemerchant/active_merchant	Rb	xetorthio/jedis	Java,HTML
activeadmin/activeadmin	Rb,SCSS	spotify/luigi	Py,JS
puppetlabs/beaker	Rb,Shell	tra/mackup	Py
boto/boto3	Py	mikel/mail	Rb
thoughtbot/bourbon	SCSS,HTML	Seldaek/monolog	PHP
bundle/bundler	Rb,HTML	omniauth/omniauth	Rb,HTML
teamcapbara/capybara	Rb	aws/opsworks-cookbooks	Rb
Codeception/Codeception	PHP,HTML	thoughtbot/paperclip	Rb
oioi/coi-services	Py,C/C++	getpelican/pelican	HTML,Py
pentaho/data-access	Java,JS	cakephp/phinx	PHP
pennersr/django-allauth	Py,HTML	propelorm/Propel2	PHP,HTML
encode/django-rest-framework	Py,HTML	puppetlabs/puppetlabs-apache	Rb
django-tastypie/django-tastypie	Py	sferik/rails_admin	JS,Rb
doorkeeper-gem/doorkeeper	Rb	reactjs/react-rails	JS,Rb
drapergem/draper	Rb,HTML	resque/resque	Rb
errbit/errbit	Rb,HTML	restsharp/RestSharp	C#
jordansissel/fpm	Rb	imperham/sidekiq	Rb,JS
ros-simulation/gazebo_ros_pkgs	C&C++	plataformatec/simple_form	Rb
ruby-grape/grape	Rb,HTML		

- 1000 commits
- 10% Bug Fixes
- 5 releases
- 30 developers
- 3 years of activity

And this resulted in 37 projects shown in Table 2. Our selected projects use many current languages such as Python, Ruby, Java, JavaScript, PHP, etc. To remove noise we ignore test cases, configuration files and static resources such as text, readme, images, etc. To achieve this first we identified source code extensions from our samples they are *py, java, rb, c, cpp, h, php, sh, cs, scss, html, scala, js, css, clj, ctp, erb, go, haml, hs and sql*. Essentially we consider only the file paths the ends with these extensions, with an additional check that the file/path names do not contain “test”, to eliminate test cases.

3.2 Data distribution

Overall this sample contains data modified in the period 2005 to 2019 by 12,361 developers in over 524,000 file entries. These modifications were made to 127,950 active branch commits (in all, 19,617,483 line insertions and 13,642,341 line deletions). On average our chosen set of projects has been active for over 7 years (see the distributions of Figure 1). The OS projects we chose for this study are publicly available. All our data is on-line so that other researchers can benefit (see github.com/ai-se/defect_perceptions).

3.3 Data Attributes

We mine the entire commit history of a project. A tuple of our main dataset contains the following attributes unique *commit_id*, *commit_time*, the author who pushed the changes *commit_author*, affected *file_path*, number of lines inserted and deleted (*insertions*, *deletions*). Importantly for every commit, we set a label *BFC* as 1, if it was pushed in an attempt to fix bug (s) or 0 if otherwise. We achieve this by scanning the commit message for if they contained any derivatives of the following stemmed words like *bug, fix, issu, error, correct, proper, deprecate, broke, optimize, patch, solve, slow,*




Figure 1: Distribution of Commits (2304), Bug Fixing (31%) Releases (60), Developers (284) and Years (8) active among 37 projects. \tilde{x} (median)

obsolete, vulnerab, debug, perf, memory, minor, wart, better, complex, break, investigat, compile, defect, inconsist, crash, problem or resol. We also identify releases of each project using Git releases/tags.

3.4 Method

Our work requires us to *measure* the following two items:

- *Effect*: Significant associations between beliefs and defect proneness.
- *Rank*: Significant differences among these associations.

3.4.1 Effect. We use Spearman’s rank correlation (a non-parametric test) to assess associations between independent variable “ F_{BX} ” and dependent variable “ F_D ” (detailed in §4). We chose Spearman like a previous defect prediction study [3] recommended to handle skewed data, further it is unaffected by transformations (log, square-root etc) on variables.

The Spearman’s rank correlation, $\rho = \frac{\text{cov}(X,Y)}{\sigma_x \sigma_y}$ between two samples X, Y (with means \bar{x} and \bar{y}), as estimated using $x_i \in X$ and $y_i \in Y$ via

$$\rho = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum_{i=1}^n (x_i - \bar{x})^2} \sqrt{\sum_{i=1}^n (y_i - \bar{y})^2}}$$

We conclude using both correlation coefficient (ρ) and its associated *p_value* in all our experiments. ρ varies from 1, i.e., ranks are identical, to -1, i.e., ranks are the opposite, where 0 indicates no correlation.

- Higher ρ value indicates more evidence for defect prone.
- Lower *p_value* indicates significant result.

3.4.2 Rank. Our population consists of a significant ρ scores. But populations may have the same median but their distribution could be very different, hence to identify significant differences or rank among two or more populations we use ScottKnott test recommended by Mittas et al. in TSE’13 paper [22]. ScottKnott is a top-down bi-clustering approach used to rank different treatments, the treatment could be beliefs, project, release or any conjunction. This method sorts a list of l treatments with l_s measurements by their median score. It then splits l into sub-lists m, n in order to maximize the expected value of differences in the observed performances before and after divisions.

For lists l, m, n of size ls, ms, ns where $l = m \cup n$, the “best” division maximizes $E(\Delta)$; i.e. the difference in the expected mean value before and after the spit:

$$E(\Delta) = \frac{ms}{ls} \text{abs}(m.\mu - l.\mu)^2 + \frac{ns}{ls} \text{abs}(n.\mu - l.\mu)^2$$

Further, to avoid “small effects” with statistically significant results we employ the conjunction of bootstrapping and A12 effect size test by Vargha and Delaney [30] for the hypothesis test H to check if m, n are truly significantly different. These techniques do not make gaussian assumptions (non-parametric).

3.5 Terminologies

In this section, we briefly introduce common labels, definitions, and thresholds that we weave while discussing the results.

We label beliefs as,

- *popular* if they have more than 50% agreement in Table 1
- *unpopular* otherwise.

3.5.1 Evidence: We use Spearman’s rank correlation (detailed in §3) ρ ranges shown below to discuss strength of the evidence. From the usage of Spearman’s ρ in this defect prediction literature [36] we derive the following ranges for $|\rho|$:

- * 0.0 to 0.39 as *no support*
- * 0.4 to 0.49 as *minimum/weak support*
- * 0.5 to 0.59 as *support*
- * 0.6 to 0.69 as *strong support*
- * 0.7 to 1.00 as *very strong support*

We acknowledge that these ranges are debatable. But importantly these ranges have lesser impact on our conclusion as we are keen on relative ranks obtained from Scott-Knott-test of §3.4.2 deduced from overall population of ρ scores.

3.5.2 Project Support Population (P_{BX}). P_{BX} is the population of significant ρ scores computed for a belief BX beliefs 1 to 10 ie, $X \in [1..10]$ for all releases R in project P . Thus,

$$P_{BX} = \{\rho(\{F_{BX}\}, \{F_D\})\}$$

To recollect, ρ is the association between:

- F_{BX} , which is the metric we collect from the Github data
- and F_D which is the defects fixed during 6 months after the release r (i.e. the defects that anyone thought were worth acting on).

We only consider significant (99% confidence level) ρ scores. This implies that some releases are not considered in the population. In other words, $|P_{BX}|$ is the number of releases in the project P , such that $|P_{BX}| \leq R$. where R is the total number of releases in a project P .

4 ASSESSING BELIEFS

In this section, we describe the construction of our correlation experiments using the underlying rationale and relevant literature attached to each of the beliefs in Table 1. Then, we discuss the role of relevant attributes scraped from our sample projects to be used as an independent variable to compute associations.



Objective:

To export P_{BX} which is a population of significant ρ scores for a belief BX , where each score is computed as a correlation between, F_{BX} & F_D and collected over all releases R in a project P . This is repeated for all 37 projects and 10 beliefs.

Where,

$r \rightarrow$ is a release in project p . We ignore first release of all projects $r > 1$.

$F \rightarrow$ is a file created or modified during the pre-release r period.

$BX \rightarrow$ denotes metric associated with beliefs 1 to 10, $X \in [1..10]$.

$F_{BX} \rightarrow$ denotes metric captured on file F

$D \rightarrow$ defects fixed 6 months after r (post-release bugs)

$F_D \rightarrow$ defects fixed 6 months after r (post-release bugs) on file F

We employ the traditional release based approach to assess our beliefs. We capture metrics in the pre-release period to find associations with post-release bugs 6 months after the release. We assess the beliefs “ BX ” for all releases in a project. We then build a population with the effect scores collected in all releases and projects to be analyzed in the §5. In each release, we compute the belief metric for each distinct file “ F ” modified in a release r where $r \in R$ (total number of releases). And defect proneness “ D ” is the number of post-release bugs. Then “ F_D ” is the number of post-release ($r + 6$ months) bugs on file “ F ”. Hence, “ F_D ” computation is common across all the beliefs whereas “ F_{BX} ” (belief metric on a file) captured during the pre-release period r is elucidated in the following sections.

In the following, we discuss the specifics of how to assess the beliefs of Table 1.

4.1 B1: Complex Code Changes

Description: “A file with a complex code change process tends to be buggy”. In 2009, A. E. Hassan [7] adopted Shannon Entropy from information theory to devise few code change models that weigh scattered modifications (complexity) of a file using its change history. The intuition behind this belief is that change scattering of a file over some periods makes it cumbersome for developers to maintain, thus making it defect prone. We use their compute History Complexity Metric (HCM^{1d}) with the decay factor (d) to assess this belief. The decay factor is used to undermine earlier modifications.

Procedure: We divide the pre-release period r into bi-weekly (14 days) periods to compute entropy HCM for a file F and this will be F_{B1} . In cases where releases have less than 14 days, r will have just two equal size periods. We then export the correlation between F_{B1} & F_D .

4.2 B2 & B10: Ownership

Description: B2: “Files changed by more developers are more buggy”. In 2010, Matsumoto et al. [18] showed human factors can be used to forecast defects. Specifically, they saw that files touched by more developers make the file prone to defects.

B10: “Files with fewer lines contributed by their owners (who contribute most changes) are bug-prone”. Here we employ the idea

of minor & major contributor introduced by Bird et al. in their work [2] where they explore various ownership related metrics. They define a minor contributor as someone who has made less than 5% changes and a major contributor as someone who has made at-least 5% or more. Thus if a file in a release has changed by a large proportion of minor contributors its prone to more defects.

Procedure: F_{B2} here is the count of distinct *commit authors* who made some change to the file in the pre-release. F_{B10} is the % of minor contributors (who contributed less than 5% code churn) for file F in the pre-release. Implying many minor contributors modifying a file makes it defect prone. We then export the correlation between F_{B2} & F_D and F_{B10} & F_D .

4.3 B3 & B9 : Code Churn

Description: B3: “A file with more added lines is more bug-prone” and B9: “A file with more removed lines is more bug-prone”. In 2005 Nagappan & Ball in [25] showed that relative code churns of files using number of added or deleted lines between subsequent versions are good indicators to forecast defects.

Procedure: We measure F_{B3} or F_{B9} for a file F by aggregating on the *number of lines added* (B3) or *number of lines deleted* (B9) only during the pre-release period r . We then export the correlation between F_{B3} & F_D and F_{B9} & F_D .

4.4 B4 & B6: Temporal

Description: B4: “Recently changed files tend to be buggy” and B6: “Recently bug-fixed files tend to be buggy”. Two of these temporal heuristics are introduced and explored by Hassan et al. [8] in their popular work, “Top 10 list for dynamic defect prediction”. The rationale here is that files modified closer to release periods may not be tested effectively and as a result, recently changed files would tend to introduce more bugs (also see [6]) in the near future (B4). Another intuition is that faults tend to arise at the same spot is a good indicator to measure defect proneness (B6).

Procedure: Attributes *commit_time* (longer value indicates more recent) and *BFC* (1 indicates bug fixing, 0 otherwise) are used to model these two beliefs. But a file can be modified (committed) multiple times in the pre-release period. Hence, we assign the maximum *commit_time* for F_{B4} , similarly we assign the maximum *commit_time* for F_{B6} ; provided *BFC* = 1 in the pre-release period. Further in F_{B6} if a file is never modified for the purpose of fixing a defect in the pre-release period, its ignored for analysis. We then export the correlation between F_{B4} & F_D and F_{B6} & F_D .

4.5 B5: Commit Churns

Description: “A commit that involves more added and removed lines is more bug-prone.”. A single commit affects one or more files with some line additions and/or deletions. Hindle et al. [10] and Hattori et al. [9] studied the nature and distribution of large commits in terms of the number of files it changed and highlighted that although large commits are rare they are unsafe to ignore. Owing to the belief statement we slightly deviate to weigh commits, purely based on the number of files. Rather we chose to weigh beliefs based on the amount of lines changes a commit pushes into the system. Since a commit can push the same amount of line changes or more with fewer files.

Procedure: Unlike other beliefs, that can be measured at file-level, a commit is a collection of files. It’s immutable, meaning it cannot be directly traced in the post-release period. Hence, we assess the impact of a commit using the files it affected. Thus, the independent variable is a commit in the pre-release period F_{B5} (represents a commit), which is an aggregate of code churn (insertions + deletions) for all the files part of the commit. Similarly, commit defect proneness is the aggregation of file defect proneness F_D in the post-release period. In other words, we collectively correlate between a file F that is part of a large commit in the pre-release period with the number of bugs introduced by F in the post-release period. We then export the correlation between F_{B5} & F_D .

4.6 B7 & B8: Something More

Description: Graves et al. in [6] found pragmatic effect among the two process-related metric B7 & B8 ie., “A file with more fixed bugs tends to be more bug-prone” and “A file with more commits is more bug-prone” through analyzing code from a telephone switching system.

Procedure: F_{B7} is the count of bug fixes on file F and F_{B8} is the count of modifications (commits) made on file F , computed during their corresponding pre-release periods. We then export the correlation between F_{B7} & F_D and F_{B8} & F_D .

5 RESULTS

In this section, we explore two of our RQ’s. For each RQ we discuss the motivation, approach, and findings.

5.1 RQ1: What beliefs are strongly supported in the data?

Motivation: Similar to the work by Devanbu et al. in [4] we compare practitioners’ agreement % with empirical evidence. Our objective here is to find beliefs that have strong support.

Approach: We compute P_{BX} for each belief BX for all the 37 projects. This results in 10 independent P_{BX} populations one for each belief. Next, we rank these populations by their median and effect difference using the Scott-Knott-test of §3.4.2. That results in Figure 2 which we compare with practitioners’ agreement % in Table 1.

Findings: Using the results in Figure 2 we report that, beliefs B10 and B5 have significantly *strong* support. Surprisingly, *unpopular* beliefs B6, B7, B8 and B10 have better support than *popular* beliefs B1, B3 & B4. For example:

- The ownership-based belief B10 with only 30% practitioner agreement shows *strong* support.
- The temporal belief B4 though with 58% practitioner agreement has relatively weaker support 0.4. B1 with the largest practitioner agreement 76% shows weak support.

Overall, only a few beliefs B2, B5 & B9 are in agreement (%) with practitioners beliefs. Of those, B5: *Large Commits*, with 57% practitioner agreement has the highest effect (0.7) among all the beliefs.

Rank	Belief	Median	IQR	
1	B9 (35%)	43	27	—●—
1	B4 (58%)	44	28	—●—
2	B3 (61%)	45	21	—●—
2	B1 (76%)	46	21	—●—
2	B6 (49%)	49	47	—●—
2	B7 (48%)	49	22	—●—
2	B2 (64%)	50	21	—●—
2	B8 (46%)	51	22	—●—
3	B10 (30%)	63	23	—●—
4	B5 (57%)	71	26	—●—

Figure 2: Scott-Knott test applied to all the 10 beliefs, that contains support ρ scores of 3,198 releases in all the 37 projects. Each row represents a population of P_{BX} scores across all 37 projects. (Higher rank indicates stronger support). IQR - Interquartile Range.

Project	P_c	$P_r\%$
Propel2	2	21
activeadmin	2	10
mackup	2	3
erbit	2	19
opsworke-cookbooks	3	4
bourbon	3	6
mail	4	13
rails_admin	4	14
doorkeeper	4	8
data-access	5	4
boto3	5	2
omniauth	6	7
bundler	6	21
capypara	6	14
active_merchant	6	4
django-allauth	6	14
monolog	7	5
Codeception	7	20
paperclip	8	9
resque	8	8
puppetlabs-apache	8	13
react-rails	8	8
RestSharp	8	26
fpm	9	14
gazebo_ros_pkgs	9	11
luigi	9	23
django-rest-framework	9	29
sidekiq	9	11
grape	10	30
pelican	10	35
jedis	10	36
draper	10	21
django-tastypie	10	30
coi-services	10	22
beaker	10	12
phoenix	10	15
simple_form	10	10

Figure 3: Coverage of beliefs P_c and its prevalence among releases P_r is shown here. Projects are sorted based on P_c . P_c is the count of beliefs that show at least a *minimum support* in the project. $P_r\%$ is the proportion of aggregated experiments that show at least a *minimum support* in the project for all the beliefs.

Result:

We found that beliefs labeled B10 & B5 in Table 1, which is believed by only (30% & 57%) of practitioners has strong support (i.e., large correlations), whereas the more popular belief B1, which is believed by 76% of practitioners, has relatively weaker support.

5.2 RQ2 : Does the same evidence appear everywhere?

Motivation: The above results show that, overall, many commonly held beliefs do not always hold across all the data. But this is not to say that *sometimes* it may be true that *some* of the beliefs of Table 1 are not true.

Devanbu et al. in [4] in their assessment found that practitioners had different opinions working for different projects, within the same organization. On similar lines, we aim to investigate irregularities of evidence among projects and releases and hence we ask:

Table 3: Small Medium Large Scott-Knott-test placed 30 treatments (10 beliefs * 3 release size) into various ranks using their support P_{BX} population. Beliefs with high support ρ are found in the bottom and less support in the top. (Higher rank indicates stronger support). IQR - Interquartile Range. Treatment labels are sub-scripted with practitioners' agreement from Table 1.

Rank	Treatment	Median (ρ)	IQR	
1	L _{B9(35%)}	33	21	—●—
1	L _{B6(49%)}	34	23	—●—
1	L _{B4(58%)}	37	22	—●—
2	L _{B1(76%)}	38	17	—●—
2	L _{B3(61%)}	38	14	—●—
2	L _{B8(46%)}	41	19	—●—
2	L _{B2(64%)}	42	18	—●—
2	L _{B7(48%)}	42	19	—●—
3	L _{B10(30%)}	52	22	—●—
3	M _{B9(35%)}	52	14	—●—
3	M _{B4(58%)}	54	16	—●—
3	M _{B7(48%)}	54	14	—●—
3	M _{B3(61%)}	54	12	—●—
3	M _{B2(64%)}	55	14	—●—
3	M _{B8(46%)}	55	13	—●—
3	M _{B1(76%)}	55	17	—●—
3	M _{B6(49%)}	55	24	—●—
4	L _{B5(57%)}	63	25	—●—
5	M _{B10(30%)}	65	15	—●—
5	S _{B1(76%)}	66	4	—●—
5	M _{B5(57%)}	69	23	—●—
5	S _{B2(64%)}	74	13	—●—
5	S _{B6(49%)}	75	22	—●—
5	S _{B9(35%)}	74	15	—●—
5	S _{B3(61%)}	76	19	—●—
5	S _{B7(48%)}	78	16	—●—
6	S _{B10(30%)}	80	15	—●—
6	S _{B4(58%)}	82	20	—●—
6	S _{B8(46%)}	80	16	—●—
6	S _{B5(57%)}	87	18	—●—

- Do projects show evidence for all the beliefs?
- Does the size of a release affect belief support?




Figure 4: Evolution of beliefs among all 37 projects. Growth % is the proportion of projects where a corresponding belief correlated with $\rho \geq 0.4$ between project release dates and empirical evidence (P_{BX}). And, Decay % is the proportion of projects where a corresponding belief correlated with $\rho \leq -0.4$ between project release dates and empirical evidence (P_{BX}). Beliefs in the plot are sorted based on Decay %.

- Do beliefs evolve as a project matures (more releases)?

a) Do projects show evidence for all the beliefs ?

Approach: For each project we measure coverage of beliefs P_C , which is simply the number of beliefs a project P shows at least a *minimum support*. In other words a project P covers a belief BX if $\tilde{x}(P_{BX}) \geq 0.4$. Then, we measure $P_r\%$ (release prevalence) which is the proportion of our aggregated correlation experiments that shows atleast a *minimum support*. To compute this we aggregate P_{BX} for all 10 beliefs for a project P . Then using this aggregated population we compute the % of $\rho \geq 0.4$.

Findings: The results in Figure 3 show that:

- Only 24% of the projects show support for all the 10 beliefs.
- Only 24% of projects show support for less than 5 beliefs (but projects that covered all 10 beliefs had their evidence in only (10 - 36%) of its releases).

b) Do beliefs fluctuate with the size of a release?

Approach: We define size of a release based on the number of files created or modified in a release. To measure fluctuations, first we group releases into three categories namely *small*, *medium* & *large* based on the number of distinct files D_F in a release. Using the distribution of D_F among releases of all the 37 projects, we find the $\tilde{x}(D_F)$ to be 18 (files). And using the Inter-Quartile range of this distribution, we place a release r of a project P in one of the following categories:

- *small*: If, $D_F < 18$ (\tilde{x})
- *medium*: If, $18 \geq D_F < Q3$ (Third quartile)
- *large*: If, $D_F \geq Q3$ (Third quartile)

Then, for each belief BX we segregate the support score population P_{BX} (computed in RQ1) into three different populations based on its release size. After grouping, this results in 30 populations (treatments) which we cluster using Scott-Knott-test§3.4.2, resulting in Table 3. Note that we group releases based on the number of files D_F rather than the median duration (21 days), since we model our experiments based on files F . This decision should not affect our experiment as we tested D_F and its corresponding release duration to be positively correlated with *strong support* of 0.6 (median); considering all releases in all the 37 projects.

Findings: We observe a clear drift in effect distribution among the three release sizes in Table 3. Notably, smaller releases have better support than larger releases. Thus, in summary, Table 3 tells us that we can reason better about smaller releases than larger releases.

That said, it is important to add that while 47% of releases in our projects are *small*, only 18% (718 releases) of them were qualified for the treatments above as the rest were insignificant ($p_value \geq 0.01$). So while our conclusions about small releases are a promising result, overall across all our data, it holds for only a very small group.

c) Do beliefs evolve as a project matures ?

Approach: To observe whether belief support tend to strengthen or decay as a project matures with more releases, we test if there is a linear relationship between belief support and its maturity. To measure this, we correlate between all release dates ($\{r_{date}\}$) of a project with its corresponding support scores P_{BX} for a belief BX .

The *Growth*(%) for a belief is the proportion of projects that shows a correlation $\rho \geq 0.4$ (indicating a positive relationship between release dates and belief support). Similarly, *Decay*(%) for a belief is the proportion of projects that shows a correlation $\rho \leq -0.4$ (indicating a negative relationship between release dates and belief support).

Findings: Figure 4 show some support for beliefs both decaying and growing. Beliefs B6 & B9 show growth in 21% of the projects. Four of the beliefs (B2, B4, B9 & B10) are decaying in more than 25% of the projects with the highest decay of 51% is observed with ownership-based belief *B10* and notably the least growth of just 2%.

That said, the major trend in Figure 4 is that beliefs tend to decay, not strengthen as a project matures. Further, we highlight that we only assessed for linearity and we don't delve into the magnitude of this effect which would remain a future work.

Result:

Only 24% of the projects show support for all the 10 beliefs. And those projects showed support for beliefs among 10% - 36% of their releases. Beliefs appear stronger in smaller than larger releases, with fluctuations like B4 (weaker in *large* releases but stronger in *small* releases). Support for beliefs tended to decay than strengthen as the project matured.

6 FUTURE DIRECTIONS

We portrayed various irregularities of beliefs in data. Krishna et al. through transfer learning techniques [14, 15] showed instabilities can be minimized by identifying a representative oracle project among ‘N’ projects. But in our opinion, as a project evolves it is more likely that such an oracle may need frequent adjustments or even replacements. Thus in a result analogous to Menzies et al. in [20], we advise focusing on factors that help to answer when & where support for beliefs holds for our future work.

Other researchers endorse our call for more reasoning about the context in software engineering. Recently Zhang et al. in [35] showed improved global predictive power by including 6 context factors (i.e., programming language, issue tracking, the total lines of code, the total number of files, the total number of commits, and the total number of developers). A comprehensive list of context factors in SE was reported by Petersen & Wohlin in [27] (2009) in our opinion is still underexplored in defect prediction. Only 3 out of 21 defect prediction related work stressed the importance of context citing Petersen & Wohlin’s context factors in the past decade (among 221 citations in total).

7 THREATS TO VALIDITY

7.1 Threats to internal validity

This paper only explored 10 of the 15 metric-related beliefs documented by Wan et al. [32] in Table 1. We found that some of the modeling decisions about how to map data into Table 1 required extensive, possibly even arcane, explanations. Objectively, we were able to answer our central question using the 10 beliefs and it would remain the same with or without exploring additional beliefs.

Like this prominent [28] and a recent [31] large scale analysis, we rely on labeling commit messages for bug fixes as our independent variable. Due to limitations in the heuristics used to generate those labels [5, 29], such labels might be misleading. To partially mitigate this issue of false negatives, we expanded our set of keywords for better coverage, which we detailed in §3. To validate if false positives impacted our results, we cross-checked whether projects with higher *Bug Fix %* were likely to show support for more beliefs.

7.2 Threats to external validity

Our work is biased by the projects and data found in our sample. This means that our conclusions could differ from other publications. For example, many of those publications discussed C and C++ projects while our sample contained many Ruby projects. Having said this, our sample sufficiently covers many popular programming languages like Python, Java, etc. All our projects are OS but Agrawal et al. in [1] showed that many open-source lessons hold for in-house. We identified releases using *git tags* which mark a boundary for readiness in the commit. Some of these changes were too small (just a few hours) even for rapid releases [17]. Hence, we only consider releases that had at least 3 distinct files changed.

7.3 Threats to statistical conclusion validity

All our conclusions are based on correlation which means the strength (or weakness) of this analysis is the same as the strength or weakness of correlation. To increase the validity of those conclusions, we only

reported significant at the 0.01 level. Also, we ignored correlations with less than 4 observations even at a significant level. As some releases have less than 3 files changed and technically it is possible to get high correlation with a zero *p-value* with just two observations, but we considered this unwise (so we deleted such conclusions).

8 CONCLUSION

At the start of a software analytics project, it is important to focus software analytics on questions of interest to the client. Therefore, it is very important to document developer beliefs, as done by Wan et al. [32]. That being said, once project data becomes available, it is just as important to update the focus in accordance with the observed effects.

In this paper, we report the source of the disconnect between practitioners’ perception of defect prediction metrics and empirical evidence. We argue sporadicity of evidence prevalence is the leading cause of the disconnect between stated beliefs (e.g. the Wan et al. [32] paper) and observed evidence. Despite irregularities, we offer few beliefs that stood out B5 (Large Commits) and B10 (Owner contribution) with good prevalence in our large scale study. Thus we encourage developers to update their beliefs when the evidence demands it. This paper also observed fluctuations of evidence in a new dimension (size of a release).

Going forward, we would advise that while it is important to report a strong belief, it is also important to take much care to report where (and when) that belief no longer holds.

ACKNOWLEDGEMENTS

This work was partially supported by NSF grant #1908762.

REFERENCES

- [1] Amritanshu Agrawal, Akond Rahman, Rahul Krishna, Alexander Sobran, and Tim Menzies. 2018. We don’t need another hero?: the impact of heroes on software development. In *Proceedings of the 40th International Conference on Software Engineering: Software Engineering in Practice*. ACM, 245–253.
- [2] Christian Bird, Nachiappan Nagappan, Brendan Murphy, Harald Gall, and Premkumar Devanbu. 2011. Don’t touch my code!: examining the effects of ownership on software quality. In *Proceedings of the 19th ACM SIGSOFT symposium and the 13th European conference on Foundations of software engineering*. ACM, 4–14.
- [3] Marco D’Ambros, Michele Lanza, and Romain Robbes. 2010. An extensive comparison of bug prediction approaches. In *2010 7th IEEE Working Conference on Mining Software Repositories (MSR 2010)*. IEEE, 31–41.
- [4] Premkumar Devanbu, Thomas Zimmermann, and Christian Bird. 2016. Belief & evidence in empirical software engineering. In *2016 IEEE/ACM 38th International Conference on Software Engineering (ICSE)*. IEEE, 108–119.
- [5] Yuanrui Fan, Xin Xia, Daniel Alencar da Costa, David Lo, Ahmed E Hassan, and Shanping Li. 2019. The Impact of Changes Misclassified by SZZ on Just-in-Time Defect Prediction. *IEEE Transactions on Software Engineering* (2019).
- [6] Todd L Graves, Alan F Karr, James S Marron, and Harvey Siy. 2000. Predicting fault incidence using software change history. *IEEE Transactions on software engineering* 26, 7 (2000), 653–661.
- [7] Ahmed E Hassan. 2009. Predicting faults using the complexity of code changes. In *Proceedings of the 31st International Conference on Software Engineering*. IEEE Computer Society, 78–88.
- [8] A. E. Hassan and R. C. Holt. 2005. The top ten list: dynamic fault prediction. In *21st IEEE International Conference on Software Maintenance (ICSM’05)*. 263–272. <https://doi.org/10.1109/ICSM.2005.91>
- [9] Lile P Hattori and Michele Lanza. 2008. On the nature of commits. In *Proceedings of the 23rd IEEE/ACM International Conference on Automated Software Engineering*. IEEE Press, III–63.
- [10] Abram Hindle, Daniel M German, and Ric Holt. 2008. What do large commits tell us?: a taxonomical study of large commits. In *Proceedings of the 2008 international working conference on Mining software repositories*. ACM, 99–108.
- [11] Thong Hoang, Hoa Khanh Dam, Yasutaka Kamei, David Lo, and Naoyasu Ubayashi. 2019. DeepJIT: an end-to-end deep learning framework for just-in-time

- defect prediction. In *Proceedings of the 16th International Conference on Mining Software Repositories*. IEEE Press, 34–45.
- [12] Magne Jørgensen and Tanja M. Gruschke. 2009. The Impact of Lessons-Learned Sessions on Effort Estimation and Uncertainty Assessments. *Software Engineering, IEEE Transactions on* 35, 3 (May-June 2009), 368–383.
 - [13] Yasutaka Kamei, Emad Shihab, Bram Adams, Ahmed E Hassan, Audris Mockus, Anand Sinha, and Naoyasu Ubayashi. 2012. A large-scale empirical study of just-in-time quality assurance. *IEEE Transactions on Software Engineering* 39, 6 (2012), 757–773.
 - [14] Rahul Krishna and Tim Menzies. 2018. Bellwethers: A baseline method for transfer learning. *IEEE Transactions on Software Engineering* (2018).
 - [15] Rahul Krishna, Tim Menzies, and Wei Fu. 2016. Too much automation? The bellwether effect and its implications for transfer learning. In *Proceedings of the 31st IEEE/ACM International Conference on Automated Software Engineering*. ACM, 122–131.
 - [16] David Lo, Nachiappan Nagappan, and Thomas Zimmermann. 2015. How practitioners perceive the relevance of software engineering research. In *Proceedings of the 2015 10th Joint Meeting on Foundations of Software Engineering*. ACM, 415–425.
 - [17] Mika V Mäntylä, Bram Adams, Foutse Khomh, Emelie Engström, and Kai Petersen. 2015. On rapid releases and software testing: a case study and a semi-systematic literature review. *Empirical Software Engineering* 20, 5 (2015), 1384–1425.
 - [18] Shinsuke Matsumoto, Yasutaka Kamei, Akito Monden, Ken-ichi Matsumoto, and Masahide Nakamura. 2010. An analysis of developer metrics for fault prediction. In *Proceedings of the 6th International Conference on Predictive Models in Software Engineering*. ACM, 18.
 - [19] Shane McIntosh and Yasutaka Kamei. 2017. Are fix-inducing changes a moving target? a longitudinal case study of just-in-time defect prediction. *IEEE Transactions on Software Engineering* 44, 5 (2017), 412–428.
 - [20] Tim Menzies, Andrew Butcher, Andrian Marcus, Thomas Zimmermann, and David Cok. 2011. Local vs. global models for effort estimation and defect prediction. In *2011 26th IEEE/ACM International Conference on Automated Software Engineering (ASE 2011)*. IEEE, 343–351.
 - [21] Tim Menzies, William Nichols, Forrest Shull, and Lucas Layman. 2017. Are delayed issues harder to resolve? Revisiting cost-to-fix of defects throughout the lifecycle. *Empirical Software Engineering* 22, 4 (2017), 1903–1935.
 - [22] N. Mittas and L. Angelis. 2013. Ranking and Clustering Software Cost Estimation Models through a Multiple Comparisons Algorithm. *IEEE Trans SE* 39, 4 (April 2013), 537–551. <https://doi.org/10.1109/TSE.2012.45>
 - [23] Nuthan Munaiah, Steven Kroh, Craig Cabrey, and Meiyappan Nagappan. 2017. Curating GitHub for engineered software projects. *Empirical Software Engineering* 22, 6 (2017), 3219–3253.
 - [24] Meiyappan Nagappan, Romain Robbes, Yasutaka Kamei, Éric Tanter, Shane McIntosh, Audris Mockus, and Ahmed E Hassan. 2015. An empirical study of goto in C code from GitHub repositories. In *Proceedings of the 2015 10th Joint Meeting on Foundations of Software Engineering*. ACM, 404–414.
 - [25] Nachiappan Nagappan and Thomas Ball. 2005. Use of relative code churn measures to predict system defect density. In *Proceedings of the 27th international conference on Software engineering*. ACM, 284–292.
 - [26] Carol Passos, Ana Paula Braun, Daniela S. Cruzes, and Manoel Mendonca. 2011. Analyzing the Impact of Beliefs in Software Project Practices. In *ESEM'11*.
 - [27] Kai Petersen and Claes Wohlin. 2009. Context in industrial software engineering research. In *2009 3rd International Symposium on Empirical Software Engineering and Measurement*. IEEE, 401–404.
 - [28] Baishakhi Ray, Daryl Posnett, Vladimir Filkov, and Premkumar Devanbu. 2014. A large scale study of programming languages and code quality in github. In *Proceedings of the 22nd ACM SIGSOFT International Symposium on Foundations of Software Engineering*. ACM, 155–165.
 - [29] Jacek Śliwerski, Thomas Zimmermann, and Andreas Zeller. 2005. When Do Changes Induce Fixes?. In *Proceedings of the 2005 International Workshop on Mining Software Repositories (MSR '05)*. ACM, New York, NY, USA, 1–5. <https://doi.org/10.1145/1082983.1083147>
 - [30] András Vargha and Harold D Delaney. 2000. A critique and improvement of the CL common language effect size statistics of McGraw and Wong. *Journal of Educational and Behavioral Statistics* 25, 2 (2000), 101–132.
 - [31] Neil Walkinshaw and Leandro Minku. 2018. Are 20% of Files Responsible for 80% of Defects?. In *Proceedings of the 12th ACM/IEEE International Symposium on Empirical Software Engineering and Measurement (ESEM '18)*. ACM, New York, NY, USA, Article 2, 10 pages. <https://doi.org/10.1145/3239235.3239244>
 - [32] Zhiyuan Wan, Xin Xia, Ahmed E Hassan, David Lo, Jianwei Yin, and Xiaohu Yang. 2018. Perceptions, expectations, and challenges in defect prediction. *IEEE Transactions on Software Engineering* (2018).
 - [33] Xin Xia, Lingfeng Bao, David Lo, Payneet Singh Kochhar, Ahmed E Hassan, and Zhenchang Xing. 2017. What do developers search for on the web? *Empirical Software Engineering* 22, 6 (2017), 3149–3185.
 - [34] Xin Xia, Zhiyuan Wan, Payneet Singh Kochhar, and David Lo. 2019. How practitioners perceive coding proficiency. In *Proceedings of the 41st International Conference on Software Engineering*. IEEE Press, 924–935.
 - [35] Feng Zhang, Audris Mockus, Iman Keivanloo, and Ying Zou. 2014. Towards building a universal defect prediction model. In *Proceedings of the 11th Working Conference on Mining Software Repositories*. ACM, 182–191.
 - [36] Thomas Zimmermann, Rahul Premraj, and Andreas Zeller. 2007. Predicting defects for eclipse. In *Third International Workshop on Predictor Models in Software Engineering (PROMISE'07: ICSE Workshops 2007)*. IEEE, 9–9.
 - [37] Weiqin Zou, David Lo, Zhenyu Chen, Xin Xia, Yang Feng, and Baowen Xu. 2018. How practitioners perceive automated bug report management techniques. *IEEE Transactions on Software Engineering* (2018).